

untitled3

September 1, 2024

1 PANDAS IN DATA SCIENCE

```
[ ]: Getting Familiar with Pandas:
```

Pandas is a powerful and widely-used library in Python for data manipulation and analysis. It provides data structures like Series and DataFrame that make it easy to work with structured data, such as tables or time series

It provides two primary data structures: Series, DataFrame

With Pandas, you can easily: Read and write data from files (CSV, Excel, SQL, etc.). Filter, sort, and modify your data. Handle missing data by filling or removing it. Group and aggregate data for summary statistics. Merge and join datasets.

1.1 HOW TO IMPORT PANDAS

```
[1]: import pandas as pd
```

1.2 Pandas Data Structures: DataFrames and Series

```
[ ]: Series:
A one-dimensional labeled array.
Can hold any data type (integers, floats, strings, etc.).
Similar to a single column in a spreadsheet.

DataFrame:
A two-dimensional, labeled data structure.
Essentially a table with rows and columns, where each column is a Series.
Can be thought of as a collection of Series objects.
```

1.3 Creating Series and DataFrames

1.3.1 creating series

```
[5]: #From a List

data = [10, 20, 30, 40, 50]
series = pd.Series(data)
print(series)
```

```
0    10
1    20
2    30
3    40
4    50
dtype: int64
```

```
[7]: # From a Dictionary:

data = {'a': 10, 'b': 20, 'c': 30}
series = pd.Series(data)
print(series)
```

```
a    10
b    20
c    30
dtype: int64
```

1.3.2 2. Creating a DataFrame:

```
[9]: #From a Dictionary:

data = {
    'Name': ['VINU', 'VARMA', 'RACHI'],
    'Age': [19, 18, 45],
    'City': ['VSP', 'VJD', 'AKP']
}
df = pd.DataFrame(data)
print(df)
```

```
   Name  Age City
0  VINU   19  VSP
1  VARMA  18  VJD
2  RACHI  45  AKP
```

```
[11]: #From a List of Lists:
```

```
data = [
    ['VINU', 19, 'VSP'],
    ['VARMA', 18, 'VJD'],
    ['RACHI', 45, 'AKP']
]
```

```
]
df = pd.DataFrame(data, columns=['Name', 'Age', 'City'])
print(df)
```

	Name	Age	City
0	VINU	19	VSP
1	VARMA	18	VJD
2	RACHI	45	AKP

1.3.3 Creating a DataFrame from a CSV File

```
[13]: # Load data from a CSV file into a DataFrame
df = pd.read_csv('data.csv')

# Display the DataFrame
print(df)
```

	Name	Age	City
0	Vinu	19	Vsp
1	varma	18	vijayawada
2	Rachi	47	Akp
3	dinesh	20	vizag

1.4 Selecting Data

Selecting Columns:

```
[15]: df = pd.DataFrame({
    'Name': ['A', 'B', 'C'],
    'Age': [25, 30, 35],
    'City': ['MUMBAI', 'HYDERABAD', 'DELHI']
})

# Select the 'Name' column
names = df['Name']
print(names)
```

0	A
1	B
2	C

Name: Name, dtype: object

```
[17]: # MULTIPLE COLUMNS

name_age = df[['Name', 'Age']]
print(name_age)
```

	Name	Age
--	------	-----

0	A	25
1	B	30
2	C	35

1.5 Selecting Rows:

1.5.1 Using .iloc[] (Index-based Selection):

```
[19]: first_row = df.iloc[0]
      print(first_row)
```

```
Name      A
Age      25
City    MUMBAI
Name: 0, dtype: object
```

```
[21]: mul_rows = df.iloc[0:2]
      print(mul_rows)
```

	Name	Age	City
0	A	25	MUMBAI
1	B	30	HYDERABAD

1.5.2 2. Filtering Rows

```
[31]: # Filter rows where 'Age' is greater than 30
      filtered_df = df[df['Age'] > 30]
      print(filtered_df)
```

	Age	City
Name		
C	35	DELHI

```
[33]: # Filter rows where 'Age' is greater than 30 and 'City' is 'Los Angeles'
      filtered_df = df[(df['Age'] > 30) & (df['City'] == 'DELHI')]
      print(filtered_df)
```

	Age	City
Name		
C	35	DELHI

1.6 Modifying Data

```
[35]: df['Salary'] = [50000, 60000, 70000]
      print(df)
```

	Age	City	Salary
Name			
A	25	MUMBAI	50000

```
B      30  HYDERABAD  60000
C      35    DELHI   70000
```

```
[61]: df = pd.read_csv('wearable_tech_sleep_quality_1.csv')
```

```
[63]: # Display the first few rows of the dataset
print("Original Data:")
print(df.head())
```

Original Data:

	Heart_Rate_Variability	Body_Temperature	Movement_During_Sleep \
0	79.934283	37.199678	1.324822
1	67.234714	36.962317	1.855481
2	82.953771	36.529815	1.207580
3	100.460597	36.176532	1.692038
4	65.316933	36.849112	0.106385

	Sleep_Duration_Hours	Sleep_Quality_Score	Caffeine_Intake_mg \
0	4.638289	1.0	107.624032
1	6.209422	1.0	104.658589
2	6.879592	10.0	0.000000
3	10.331531	1.0	116.990981
4	8.334830	1.0	223.282908

	Stress_Level	Bedtime_Consistency	Light_Exposure_hours
0	2.771837	0.657037	7.933949
1	3.738138	0.144464	6.992699
2	3.115880	0.642949	7.655250
3	3.904008	0.453255	9.429463
4	4.571699	0.641492	10.555713

```
[65]: # Step 2: Handle Missing Data
# 2.1: Check for missing data
print("\nMissing Data Check:")
print(df.isnull().sum())
```

```
Missing Data Check:
Heart_Rate_Variability    0
Body_Temperature          0
Movement_During_Sleep    0
Sleep_Duration_Hours      0
Sleep_Quality_Score       0
Caffeine_Intake_mg        0
Stress_Level              0
Bedtime_Consistency       0
Light_Exposure_hours      0
dtype: int64
```

```
[87]: # Handle missing values
df_filled = df.fillna(df.mean(numeric_only=True))
```

```
[89]: print("\nDataFrame after handling missing values:")
print(df_filled.head())
```

DataFrame after handling missing values:

	Heart_Rate_Variability	Body_Temperature	Movement_During_Sleep \
0	79.934283	37.199678	1.324822
1	67.234714	36.962317	1.855481
2	82.953771	36.529815	1.207580
3	100.460597	36.176532	1.692038
4	65.316933	36.849112	0.106385

	Sleep_Duration_Hours	Sleep_Quality_Score	Caffeine_Intake_mg \
0	4.638289	1.0	107.624032
1	6.209422	1.0	104.658589
2	6.879592	10.0	0.000000
3	10.331531	1.0	116.990981
4	8.334830	1.0	223.282908

	Stress_Level	Bedtime_Consistency	Light_Exposure_hours
0	2.771837	0.657037	7.933949
1	3.738138	0.144464	6.992699
2	3.115880	0.642949	7.655250
3	3.904008	0.453255	9.429463
4	4.571699	0.641492	10.555713

```
[91]: # Check columns in the DataFrame
print("\nColumns in DataFrame:")
print(df_filled.columns)

# Check the first few rows to confirm 'Date' column
print("\nFirst few rows of DataFrame:")
print(df_filled.head())
```

Columns in DataFrame:

```
Index(['Heart_Rate_Variability', 'Body_Temperature', 'Movement_During_Sleep',
      'Sleep_Duration_Hours', 'Sleep_Quality_Score', 'Caffeine_Intake_mg',
      'Stress_Level', 'Bedtime_Consistency', 'Light_Exposure_hours'],
      dtype='object')
```

First few rows of DataFrame:

	Heart_Rate_Variability	Body_Temperature	Movement_During_Sleep \
0	79.934283	37.199678	1.324822
1	67.234714	36.962317	1.855481

2	82.953771	36.529815	1.207580
3	100.460597	36.176532	1.692038
4	65.316933	36.849112	0.106385

	Sleep_Duration_Hours	Sleep_Quality_Score	Caffeine_Intake_mg \
0	4.638289	1.0	107.624032
1	6.209422	1.0	104.658589
2	6.879592	10.0	0.000000
3	10.331531	1.0	116.990981
4	8.334830	1.0	223.282908

	Stress_Level	Bedtime_Consistency	Light_Exposure_hours
0	2.771837	0.657037	7.933949
1	3.738138	0.144464	6.992699
2	3.115880	0.642949	7.655250
3	3.904008	0.453255	9.429463
4	4.571699	0.641492	10.555713

1.7 DATA TRANSFORMATION

```
[93]: # For example, let's add a new column that calculates sleep efficiency
      ↪(assuming we have sleep_duration and awake_duration columns)
      if 'sleep_duration' in df_filled.columns and 'awake_duration' in df_filled.
      ↪columns:
          df_filled['sleep_efficiency'] = (df_filled['sleep_duration'] -
      ↪df_filled['awake_duration']) / df_filled['sleep_duration']
          df_filled['sleep_efficiency'] = df_filled['sleep_efficiency'].apply(lambda
      ↪x: x * 100) # convert to percentage
```

```
[95]: print("\nData After Transformation:")
      print(df_filled.head())
```

Data After Transformation:

	Heart_Rate_Variability	Body_Temperature	Movement_During_Sleep \
0	79.934283	37.199678	1.324822
1	67.234714	36.962317	1.855481
2	82.953771	36.529815	1.207580
3	100.460597	36.176532	1.692038
4	65.316933	36.849112	0.106385

	Sleep_Duration_Hours	Sleep_Quality_Score	Caffeine_Intake_mg \
0	4.638289	1.0	107.624032
1	6.209422	1.0	104.658589
2	6.879592	10.0	0.000000
3	10.331531	1.0	116.990981
4	8.334830	1.0	223.282908

	Stress_Level	Bedtime_Consistency	Light_Exposure_hours
0	2.771837	0.657037	7.933949
1	3.738138	0.144464	6.992699
2	3.115880	0.642949	7.655250
3	3.904008	0.453255	9.429463
4	4.571699	0.641492	10.555713

Pandas functions to clean and preprocess data, such as handling missing values, removing duplicates, and data type conversions.

```
[ ]: Handling Missing Values
      Removing Duplicates
      Data Type Conversions
```

```
[97]: # Display initial data summary
print("Initial Data Summary:")
print(df.info())
print("\nInitial Data Sample:")
print(df.head())
```

Initial Data Summary:

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 1000 entries, 0 to 999

Data columns (total 9 columns):

#	Column	Non-Null Count	Dtype
0	Heart_Rate_Variability	1000 non-null	float64
1	Body_Temperature	1000 non-null	float64
2	Movement_During_Sleep	1000 non-null	float64
3	Sleep_Duration_Hours	1000 non-null	float64
4	Sleep_Quality_Score	1000 non-null	float64
5	Caffeine_Intake_mg	1000 non-null	float64
6	Stress_Level	1000 non-null	float64
7	Bedtime_Consistency	1000 non-null	float64
8	Light_Exposure_hours	1000 non-null	float64

dtypes: float64(9)

memory usage: 70.4 KB

None

Initial Data Sample:

	Heart_Rate_Variability	Body_Temperature	Movement_During_Sleep \
0	79.934283	37.199678	1.324822
1	67.234714	36.962317	1.855481
2	82.953771	36.529815	1.207580
3	100.460597	36.176532	1.692038
4	65.316933	36.849112	0.106385

	Sleep_Duration_Hours	Sleep_Quality_Score	Caffeine_Intake_mg \
--	----------------------	---------------------	----------------------

0	4.638289	1.0	107.624032
1	6.209422	1.0	104.658589
2	6.879592	10.0	0.000000
3	10.331531	1.0	116.990981
4	8.334830	1.0	223.282908

	Stress_Level	Bedtime_Consistency	Light_Exposure_hours
0	2.771837	0.657037	7.933949
1	3.738138	0.144464	6.992699
2	3.115880	0.642949	7.655250
3	3.904008	0.453255	9.429463
4	4.571699	0.641492	10.555713

1.7.1 REMOVING DUPLICATES

```
[99]: # 2. Removing Duplicates

# Check for duplicate rows
duplicates = df_filled.duplicated().sum()
print(f"\nNumber of Duplicate Rows: {duplicates}")

# Remove duplicate rows
df_cleaned = df_filled.drop_duplicates()

print("\nData After Removing Duplicates:")
print(df_cleaned.head())
```

Number of Duplicate Rows: 0

Data After Removing Duplicates:

	Heart_Rate_Variability	Body_Temperature	Movement_During_Sleep \
0	79.934283	37.199678	1.324822
1	67.234714	36.962317	1.855481
2	82.953771	36.529815	1.207580
3	100.460597	36.176532	1.692038
4	65.316933	36.849112	0.106385

	Sleep_Duration_Hours	Sleep_Quality_Score	Caffeine_Intake_mg \
0	4.638289	1.0	107.624032
1	6.209422	1.0	104.658589
2	6.879592	10.0	0.000000
3	10.331531	1.0	116.990981
4	8.334830	1.0	223.282908

	Stress_Level	Bedtime_Consistency	Light_Exposure_hours
0	2.771837	0.657037	7.933949
1	3.738138	0.144464	6.992699

2	3.115880	0.642949	7.655250
3	3.904008	0.453255	9.429463
4	4.571699	0.641492	10.555713

```
[ ]: Pandas to perform data analysis, including generating summary statistics,
grouping data, and applying aggregate functions.
```

```
[ ]: Generating Summary Statistics
Grouping Data
Applying Aggregate Functions
```

2 1. Generating Summary Statistics

```
[101]: # Summary statistics for numeric columns
print("Summary Statistics for Numeric Columns:")
print(df.describe())
```

Summary Statistics for Numeric Columns:

	Heart_Rate_Variability	Body_Temperature	Movement_During_Sleep \
count	1000.000000	1000.000000	1000.000000
mean	70.386641	36.535418	2.005834
std	19.584319	0.498727	0.983454
min	5.174653	35.029806	-1.019512
25%	57.048194	36.196879	1.352000
50%	70.506012	36.531539	1.999749
75%	82.958878	36.864441	2.660915
max	147.054630	38.096554	5.926238

	Sleep_Duration_Hours	Sleep_Quality_Score	Caffeine_Intake_mg \
count	1000.000000	1000.000000	1000.000000
mean	7.471921	2.592946	148.260148
std	1.540699	2.979500	94.031760
min	3.105827	1.000000	0.000000
25%	6.393869	1.000000	80.630719
50%	7.500277	1.000000	145.717293
75%	8.500418	2.537789	211.244685
max	12.364639	10.000000	400.000000

	Stress_Level	Bedtime_Consistency	Light_Exposure_hours
count	1000.000000	1000.000000	1000.000000
mean	4.940956	0.504222	8.036684
std	2.032708	0.204137	2.023371
min	0.000000	0.000000	0.326689
25%	3.489725	0.361569	6.726291
50%	4.890507	0.500996	8.038248
75%	6.399490	0.644680	9.354408
max	10.000000	1.000000	14.754766

3 2. Grouping Data

```
[121]: if 'Body_Temperature' in df.columns:
        grouped_data = df.groupby('Body_Temperature').mean()
        print("\nGrouped Data (Mean of Numeric Columns) by 'Body_Temperature':")
        print(grouped_data)
```

Grouped Data (Mean of Numeric Columns) by 'Body_Temperature':

	Heart_Rate_Variability	Movement_During_Sleep \
Body_Temperature		
35.029806	56.378967	1.931366
35.039325	60.446851	1.542698
35.051872	61.587094	1.097948
35.063869	58.462163	2.054934
35.075729	66.286820	1.862628
...
37.794782	83.924127	2.902277
37.800842	80.080930	1.284240
37.822172	78.800289	2.876047
38.068874	79.135064	1.763445
38.096554	72.584424	2.668340

	Sleep_Duration_Hours	Sleep_Quality_Score \
Body_Temperature		
35.029806	5.846595	1.000000
35.039325	8.406181	1.000000
35.051872	8.783333	1.000000
35.063869	7.259460	1.000000
35.075729	7.000709	1.000000
...
37.794782	6.853312	1.000000
37.800842	6.659089	6.832218
37.822172	7.590129	1.000000
38.068874	7.140747	1.000000
38.096554	8.601956	10.000000

	Caffeine_Intake_mg	Stress_Level	Bedtime_Consistency \
Body_Temperature			
35.029806	150.075551	4.896916	0.628938
35.039325	124.604829	2.644477	0.621590
35.051872	200.168505	5.461222	0.474519
35.063869	245.246197	2.953494	0.437812
35.075729	254.691179	6.444894	0.619508
...
37.794782	243.864002	5.401397	0.288193
37.800842	39.188754	3.761754	0.704325

37.822172	317.344189	6.757928	0.462061
38.068874	210.258027	8.612178	0.931824
38.096554	0.000000	4.996572	0.283999

	Light_Exposure_hours
Body_Temperature	
35.029806	10.958964
35.039325	4.727194
35.051872	6.642039
35.063869	9.292259
35.075729	8.301205
...	...
37.794782	5.749604
37.800842	6.750331
37.822172	9.253754
38.068874	7.285170
38.096554	6.043954

[1000 rows x 8 columns]

3.1 AGGREGATING DATA

```
[123]: if 'Movement_During_Sleep' in df.columns:
        aggregated_data = df.groupby('Body_Temperature').agg({
            'Movement_During_Sleep': ['mean', 'sum', 'max', 'min']
        })
        print("\nAggregated Data (Mean, Sum, Max, Min) of 'Movement_During_Sleep'
        by 'Body_Temperature':")
        print(aggregated_data)
```

Aggregated Data (Mean, Sum, Max, Min) of 'Movement_During_Sleep' by
'Body_Temperature':

	Movement_During_Sleep			
	mean	sum	max	min
Body_Temperature				
35.029806	1.931366	1.931366	1.931366	1.931366
35.039325	1.542698	1.542698	1.542698	1.542698
35.051872	1.097948	1.097948	1.097948	1.097948
35.063869	2.054934	2.054934	2.054934	2.054934
35.075729	1.862628	1.862628	1.862628	1.862628
...
37.794782	2.902277	2.902277	2.902277	2.902277
37.800842	1.284240	1.284240	1.284240	1.284240
37.822172	2.876047	2.876047	2.876047	2.876047
38.068874	1.763445	1.763445	1.763445	1.763445
38.096554	2.668340	2.668340	2.668340	2.668340

[1000 rows x 4 columns]

3.2 advanced data manipulation techniques like merging, joining, and concatenating DataFrames. for above data set

```
[125]: # Create sample DataFrames for demonstration
data1 = {
    'id': [1, 2, 3, 4],
    'sleep_duration': [7.5, 6.2, 8.0, 5.5],
    'steps': [10000, 8000, 12000, 6000]
}
data2 = {
    'id': [3, 4, 5, 6],
    'calories_burned': [300, 250, 350, 200],
    'heart_rate': [70, 65, 75, 60]
}

df1 = pd.DataFrame(data1)
df2 = pd.DataFrame(data2)
```

3.3 1. Merging DataFrames

```
[127]: merged_df = pd.merge(df1, df2, on='id', how='inner') # Inner join
print("Merged DataFrame (inner join on 'id'):")
print(merged_df)
```

Merged DataFrame (inner join on 'id'):

	id	sleep_duration	steps	calories_burned	heart_rate
0	3	8.0	12000	300	70
1	4	5.5	6000	250	65

3.4 2. Joining DataFrames

```
[129]: # Set 'id' as the index for joining
df1.set_index('id', inplace=True)
df2.set_index('id', inplace=True)

# Join df1 with df2 on the index
joined_df = df1.join(df2, how='left') # Left join
print("\nJoined DataFrame (left join on index):")
print(joined_df)
```

Joined DataFrame (left join on index):

	sleep_duration	steps	calories_burned	heart_rate
id				

1	7.5	10000	NaN	NaN
2	6.2	8000	NaN	NaN
3	8.0	12000	300.0	70.0
4	5.5	6000	250.0	65.0

3.5 3. Concatenating DataFrames

```
[131]: # Concatenate df1 and df2 horizontally (adding columns)
concat_horiz_df = pd.concat([df1, df2], axis=1)
print("\nConcatenated DataFrame (horizontally):")
print(concat_horiz_df)
```

Concatenated DataFrame (horizontally):

	sleep_duration	steps	calories_burned	heart_rate
id				
1	7.5	10000.0	NaN	NaN
2	6.2	8000.0	NaN	NaN
3	8.0	12000.0	300.0	70.0
4	5.5	6000.0	250.0	65.0
5	NaN	NaN	350.0	75.0
6	NaN	NaN	200.0	60.0

4 Application in Data Science:

Pandas is a powerful library for data manipulation and analysis in Python, offering several advantages over traditional Python data structures for handling and analyzing data. Here's how the use of Pandas in the provided program can help a data science professional, and why it's beneficial compared to using native Python data structures:

4.1 Advantages of Using Pandas:

Efficient Data Handling: DataFrames and Series: Pandas introduces the DataFrame and Series objects, which are optimized for handling large datasets efficiently. These structures allow for easy manipulation of data with labels and are designed to handle data operations that would be cumbersome with lists or dictionaries.

Built-in Functions for Data Cleaning: Handling Missing Data: Pandas provides convenient methods like `fillna()`, `dropna()`, and `isnull()` to handle missing values. **Removing Duplicates:** With methods such as `drop_duplicates()`, Pandas simplifies the process of identifying and removing duplicate records.

Advanced Data Manipulation: Merging and Joining: Pandas supports sophisticated data merging and joining operations through functions like `merge()` and `join()`. **Concatenation:** Using `concat()`, Pandas can efficiently concatenate DataFrames either vertically or horizontally. This capability is essential for combining multiple datasets and performing batch processing.

Data Aggregation and Analysis: GroupBy and Aggregation: Pandas allows for easy grouping of data using `groupby()` and performing aggregate functions such as `mean()`, `sum()`, and `count()`.

This functionality enables data scientists to perform complex data analysis and summarization with concise code, rather than implementing these operations from scratch.

Flexible Data Transformation: **Data Type Conversion:** With methods like `astype()` and `pd.to_datetime()`, Pandas facilitates the conversion of data types, making it easier to ensure that data is in the correct format for analysis. **Summary Statistics:** The `describe()` method provides a quick way to generate summary statistics, helping data scientists to understand the distribution and characteristics of their data at a glance.

4.1.1 real-world examples where Pandas is essential, such as in data cleaning,

exploratory data analysis (EDA)

Pandas is a cornerstone of data science and is essential in many real-world scenarios. Here are some examples where Pandas proves invaluable, particularly in data cleaning and exploratory data analysis (EDA):

1. **Data Cleaning Example: Financial Transactions Data Scenario:** A financial institution receives daily transaction data from multiple sources, including internal systems and external partners. The data includes transaction amounts, timestamps, and account information. **Challenges:** **Missing Values:** Some transactions may have missing amounts or incomplete timestamps. **Inconsistent Data Formats:** Date and time formats might vary across different sources. **Duplicate Records:** Duplicate transactions could be recorded due to system errors. **How Pandas Helps:** **Handling Missing Values:** Pandas' `fillna()` and `dropna()` functions can be used to fill in or remove missing data. For example, `df.fillna(df['amount'].mean())` fills missing amounts with the average transaction amount.
2. **Exploratory Data Analysis (EDA) Example: E-Commerce Sales Analysis Scenario:** An e-commerce company wants to understand its sales performance over the past year. The dataset includes sales transactions, product categories, customer information, and timestamps. **Challenges:** **Summary Statistics:** Understanding the central tendencies and dispersion of sales data. **Group Analysis:** Analyzing sales performance by product category or region. **Trend Analysis:** Identifying trends and patterns over time. **How Pandas Helps:** **Generating Summary Statistics:** Use `df.describe()` to obtain summary statistics like mean, median, and standard deviation for numeric columns, which helps in understanding overall sales performance. **Grouping and Aggregation:** `df.groupby('category').agg({'sales_amount': ['mean', 'sum']})` aggregates sales data by product category, showing the average and total sales per category. **Trend Analysis:** Use `df.groupby(df['date'].dt.to_period('M')).sum()` to analyze monthly sales trends, helping to identify seasonal patterns and growth over time.